

Artificial Intelligence

6. CSP

Shashi Prabh

School of Engineering and Applied Science
Ahmedabad University

Contents

Goal: use factored representation of agents to solve problems.

Topics

- Constraint Satisfaction Problem
- Constraint Propagation
- Backtracking Search
- Local Search

Constraint Satisfaction Problems (CSP)

- We consider factored representation of states
 - A state is a set of variables
 - A problem solution is an assignment of values to the state variables where all the constraints on the variables are satisfied

Constraint Satisfaction Problems (CSP)

- Why CSP?
 - CSP is a natural formulation in many problems
 - Scheduling, planning, resource allocation, temporal models, control etc.
 - Significant reduction of search space, availability of fast solvers
 - Insight into the problem structure can be used for search speed-up
 - Some intractable atomic search-space problems can be quickly solved as CSP formulation
 - Actions and transition model can be deduced from the

Constraint Satisfaction Problems (CSP)

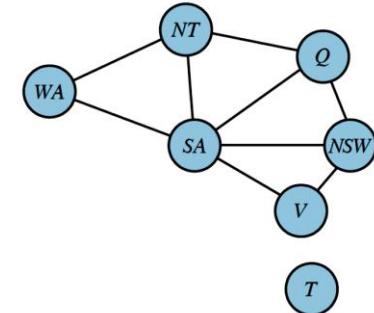
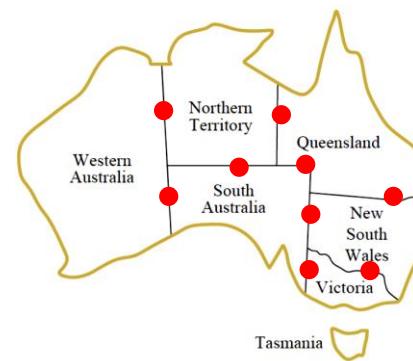
- A CSP consists of three components (X , D , C):
- **Variables** $X = \{x_1, x_2, \dots, x_n\}$
- **Domains** $D = \{D_1, D_2, \dots, D_n\}$
- **Constraints** $C = \{c_1, c_2, \dots, c_m\}$
 - Domain D_i consists of the set of **allowable values** $\{v_1, \dots, v_k\}$ for each x_i
 - $\{T, F\}$ for a Boolean variable
 - Constraint c_j consists of a pair **\langle scope, relation \rangle**
 - $\langle(x_1, x_2), x_1 \neq x_2 \rangle$ or just $x_1 \neq x_2$

Constraint Satisfaction Problems (CSP)

- Goal: Assign values to the variables from their respective domains such that all the constraints are satisfied
 - An assignment that does not violate any constraint is called **consistent** or **legal assignment**
 - A solution to a CSP is a complete and consistent assignment
 - Solving a CSP is NP-complete in general
- Types of constraints:
 - Unary, binary, global
 - $\langle (x_1, x_2) \mid x_1 \neq x_2 \rangle$, $\langle (x_1, x_2) \mid x_1 \neq x_2 \rangle$, AllDiff, AtMost, AtLeast

Map coloring

- $X = \{W, N, S, Q, NSW, V, T\}$
- $D = \{r, g, b\}$
- $C = \{ W \neq N, S \neq N, Q \neq N, W \neq S, S \neq Q, \text{etc} \}$
 - $W \neq N$ means $\{(r, g), (r, g), (g, r), (g, b), (b, r), (b, g)\}$
 - Note the reduced search space due to the constraints: 2^5 vs 3^5
 - **Can you find one solution?**
- In a CSP **constraint graph**, two variables are connected by an edge if there is a constraint that involves both



Job-Shop Scheduling – Car Assembly

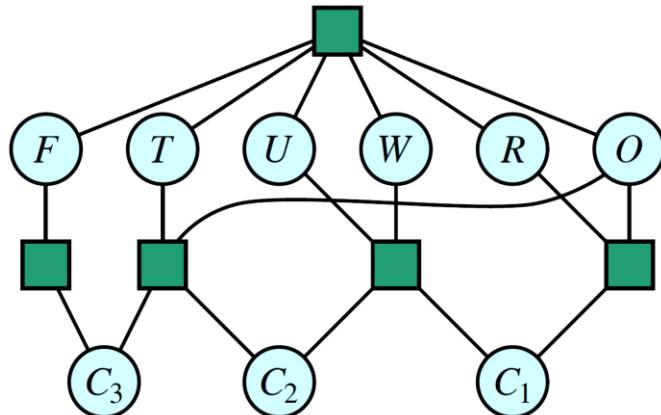
- X is the set of tasks

$\{Axe_F, Axe_B, Wheel_{RF}, Wheel_{LF}, Wheel_{RB}, Wheel_{LB}, Nuts_{RF}, Nuts_{LF}, Nuts_{RB}, Nuts_{LB}, Cap_{RF}, Cap_{LF}, Cap_{RB}, Cap_{LB}, Inspect\}$

- Values are the start times of tasks: $D_i = \{0, 1, \dots, 30\}$
- Constraints: precedence constraints, completion times
 - If T_i precedes T_j , $T_i + D_i \leq T_j$
$$Axe_F + 10 \leq Wheel_{RF}; \quad Axe_F + 10 \leq Wheel_{LF}; \\ Axe_B + 10 \leq Wheel_{RB}; \quad Axe_B + 10 \leq Wheel_{LB}.$$
 - Disjunctions: Axle installations must not overlap in time
$$(Axe_F + 10 \leq Axe_B) \quad \text{or} \quad (Axe_B + 10 \leq Axe_F)$$
 - Exercise: CSP formulation of 8-Queens problem

Cryptarithmetic Puzzles

$$\begin{array}{r} T \quad W \quad O \\ + \quad T \quad W \quad O \\ \hline F \quad O \quad U \quad R \end{array}$$



Constraint hypergraph

- Constraints: AllDiff (F, T, U, W, R, O), F ≠ 0 and

$$O + O = R + 10 \cdot C_1$$

$$C_1 + W + W = U + 10 \cdot C_2$$

$$C_2 + T + T = O + 10 \cdot C_3$$

$$C_3 = F,$$

1. In the context of a Constraint Satisfaction Problem (CSP), what are the three main components?
 - A. Variables, Functions, and Relations
 - B. Nodes, Arcs, and Paths
 - C. States, Transitions, and Goals
 - D. Variables, Domains, and Constraints
2. Which of the following is an example of a binary constraint?
 - A. The sum of variables A, B, and C must be less than 10.
 - B. The value of variable A must be different from the value of variable B.
 - C. The colors of all bordering regions must be different.
 - D. The value of variable A must be even.

1. In the context of a Constraint Satisfaction Problem (CSP), what are the three main components?
 - A. Variables, Functions, and Relations
 - B. Nodes, Arcs, and Paths
 - C. States, Transitions, and Goals
 - D. Variables, Domains, and Constraints
2. Which of the following is an example of a binary constraint?
 - A. The sum of variables A, B, and C must be less than 10.
 - B. The value of variable A must be different from the value of variable B.
 - C. The colors of all bordering regions must be different.
 - D. The value of variable A must be even.

3. In the N-queens problem, what do the variables, domains, and constraints represent?
- A. Variables are the queens, domains are the squares on the board, and constraints are that no two queens can attack each other.
 - B. Variables are the rows and columns, domains are whether a queen is present, and constraints are that a queen must be in every row.
 - C. Variables are the rows, domains are the columns, and constraints are that queens cannot be in the same row.
 - D. Variables are the columns, domains are the rows, and constraints are that no two queens can share a row, column, or diagonal.
4. Which of the following problems can be formulated as a CSP?
- A. Generating a creative story given a set of keywords.
 - B. Calculating the final grade for a student.
 - C. Determining a weekly schedule for classes without any conflicts.
 - D. Finding the shortest path between two cities on a map.

1. In the N-queens problem, what do the variables, domains, and constraints represent?
 - A. Variables are the queens, domains are the squares on the board, and constraints are that no two queens can attack each other.
 - B. Variables are the rows and columns, domains are whether a queen is present, and constraints are that a queen must be in every row.
 - C. Variables are the rows, domains are the columns, and constraints are that queens cannot be in the same row.
 - D. Variables are the columns, domains are the rows, and constraints are that no two queens can share a row, column, or diagonal.
2. Which of the following problems can be formulated as a CSP?
 - A. Generating a creative story given a set of keywords.
 - B. Calculating the final grade for a student.
 - C. Determining a weekly schedule for classes without any conflicts.
 - D. Finding the shortest path between two cities on a map.

Constraint Propagation

- A CSP algorithm can generate successors as new assignments
- Constraint propagation is an alternative where using constraints the number of legal values is reduced
- Used along-with search and/or as a preprocessing step

Constraint Propagation

- Local consistency shrinks the search space by eliminating the inconsistent assignments
- Types of local consistency
 - Node consistency
 - Arc consistency
 - Path and K-Consistency
- Global constraints, bounds propagation

Node Consistency

- A node in the constraint graph is node-consistent if all the values in the variable's domain satisfy the variable's unary constraints.
- Example: consider a unary constraint $SA \neq \{\text{green}\}$
 - The variable SA with initial domain {red, green, blue} can be made node consistent by eliminating green from its domain, leaving SA with the reduced domain {red, blue}.
- A graph is node-consistent if every variable in the graph is node-consistent.
 - One can just eliminate domain values inconsistent with unary constraints.

Arc Consistency

- A variable is arc-consistent if for every value in its domain, there is some value in the domains of all the variables connected by a binary constraint.
 - Example: consider the constraint $Y = X^2$, $D_X = \mathbb{N}$, $D_Y = \{ 1, 4, 9\}$
 - X is made arc-consistent with Y by restricting $D_X = \{ 1, 2, 3\}$
 - However, arc-consistency is ineffective in the map coloring example
- **AC-3** is a widely used arc-consistency algorithm

AC-3 (Mackworth, 1977)

function AC-3(*csp*) **returns** false if an inconsistency is found and true otherwise
queue \leftarrow a queue of arcs, initially all the arcs in *csp*

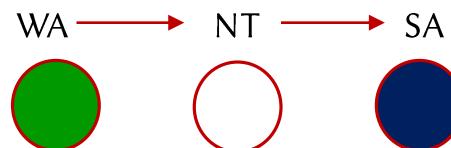
while *queue* is not empty **do**
 $(X_i, X_j) \leftarrow \text{POP}(\text{queue})$
 if REVISE(*csp*, *X_i*, *X_j*) **then**
 if size of *D_i* = 0 **then return** false
 for each *X_k* **in** *X_i.NEIGHBORS - {X_j}* **do**
 add (*X_k*, *X_i*) to *queue*
return true

function REVISE(*csp*, *X_i*, *X_j*) **returns** true iff we revise the domain of *X_i*
revised \leftarrow false
for each *x* **in** *D_i* **do**
 if no value *y* in *D_j* allows (*x,y*) to satisfy the constraint between *X_i* and *X_j* **then**
 delete *x* from *D_i*
 revised \leftarrow true
return *revised*

- Initially, each binary constraint inserts two arcs
- X_i is being made consistent with X_j
- $O(c d^3)$ worst-case complexity

Path Consistency

- AC does not help with map coloring
 - Does not object to 2-coloring the map
- $\{X_i, X_j\}$ is path-consistent with respect to a third variable X_m if, for every assignment $\{X_i = a, X_j = b\}$ consistent with the constraints on $\{X_i, X_j\}$, there is an assignment to X_m that satisfies the constraints on $\{X_i, X_m\}$ and $\{X_m, X_j\}$.
 - Refers to the overall consistency of the path $X_i \rightarrow X_m \rightarrow X_j$
 - Can infer that no valid 2-coloring of the Australia map exists



K-Consistency

- A CSP is **k-consistent** if, for any set of **k-1** variables and for any consistent assignment to those variables, a consistent value can always be assigned to any k^{th} variable
 - 1-consistency says that, given the empty set, we can make any set of one variable consistent: this is what we called node consistency
 - 2-consistency is the same as arc consistency
 - 3-consistency (binary constraints) is the same as path consistency

K-Consistency

- A CSP is **strongly k-consistent** if it is k-consistent and is also $(k-1)$ -consistent, $(k-2)$ -consistent, . . . , 1-consistent
- Why?
 - Can design a greedy algorithm
 - **CSP is NP-complete**
 - K-consistency requires exponential time and space

Global constraints

- A global constraint involves an arbitrary number of variables. It is more efficient to handle these by special-purpose algorithms
 - **AllDiff**: if m variables are involved in an AllDiff constraint, and if n possible distinct values altogether are available, then the constraint cannot be satisfied if $m > n$
 - **Atmost**: resource constraint
 - Example: no more than 10 personnel are scheduled in total
 - We can detect an inconsistency simply by checking the sum of the **minimum** values of the current domains

Global constraints

- **Bounds propagation:** For problems with large integer domains it is usually not efficient to represent the domain of each variable as a large set of integers.
 - Domains can be represented by upper and lower bounds and managed by bounds propagation

Global constraints

- Example:
 - Consider two flights, F1 and F2, for which the planes have capacities 165 and 385, respectively
 - The initial domains for the numbers of passengers are then $D_1 = [0, 165]$ and $D_2 = [0, 385]$
 - The additional constraint that **the two flights together must carry 450 people** can be handled by propagating bounds constraints as $D_1 = [65, 165]$ and $D_2 = [285, 385]$

Sudoku

	1	2	3	4	5	6	7	8	9
A			3	2		6			
B	9		3		5				1
C			1	8		6	4		
D			8	1		2	9		
E	7								8
F			6	7		8	2		
G			2	6		9	5		
H	8		2		3				9
I			5		1	3			

	1	2	3	4	5	6	7	8	9
A	4	8	3	9	2	1	6	5	7
B	9	6	7	3	4	5	8	2	1
C	2	5	1	8	7	6	4	9	3
D	5	4	8	1	3	2	9	7	6
E	7	2	9	5	6	4	1	3	8
F	1	3	6	7	9	8	2	4	5
G	3	7	2	6	8	9	5	1	4
H	8	1	4	2	5	3	7	6	9
I	6	9	5	4	1	7	3	8	2

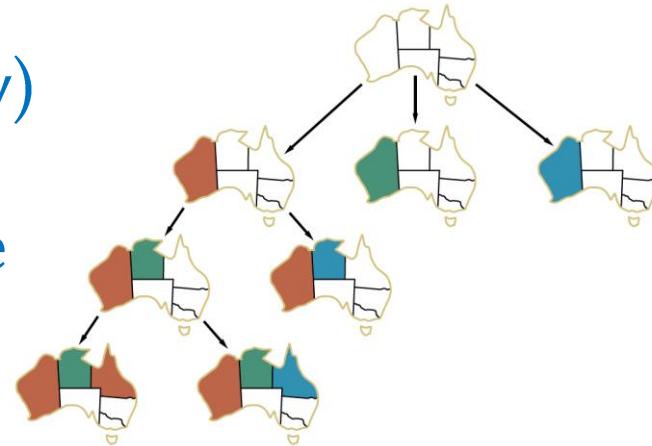
Exercise: Write CSP formulation!

Backtracking Search

- DFS at heart but with significant differences
- Search for solution is needed when after constraint propagation there exist variables with ≥ 1 values
- For a CSP with n variables of domain size d , a pure DFS results in a search tree with $n! d^n$ leaf nodes at depth n
 - The branching factor at the top would be nd , at the next level $(n-1)d$ and so on
- The order of assignments does not matter
 - There are only d^n possible assignments!

Backtracking Search

- Backtracking search progresses via a recursive call
- An unassigned variable is (repeatedly) chosen, a value is assigned and the search progresses to another variable
 - If the search succeeds, the solution is returned
 - If the search fails, the assignment is restored to the previous state, and the next value is tried



Backtracking Search

```
function BACKTRACKING-SEARCH(csp) returns a solution or failure
    return BACKTRACK(csp, {})

function BACKTRACK(csp, assignment) returns a solution or failure
    if assignment is complete then return assignment
    var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(csp, assignment)
    for each value in ORDER-DOMAIN-VALUES(csp, var, assignment) do
        if value is consistent with assignment then
            add  $\{var = value\}$  to assignment
            inferences  $\leftarrow$  INFERENCE(csp, var, assignment)
            if inferences  $\neq$  failure then
                add inferences to csp
                result  $\leftarrow$  BACKTRACK(csp, assignment)
                if result  $\neq$  failure then return result
                remove inferences from csp
                remove  $\{var = value\}$  from assignment
    return failure
```

Improving Backtracking Search

- Backtracking search can be improved using **domain-independent heuristics** that take advantage of the factored representation of states
- Variable and value ordering heuristics
 - **Minimum-remaining-values heuristic (MRV)**
 - Start with F in cryptarithmetric puzzle
 - **Degree heuristic – largest first**
 - Start with SA in Australia map
 - **Least constraining value first heuristic (LCV)**
 - Values that rule out the fewest choices first

Forward Checking

- **Forward Checking:** Check for arc consistency upon a variable assignment
 - Upon assignment to X, make each unassigned variable Y that is connected to X by a constraint, arc-consistent with X
 - After assigning V =blue, forward checking finds domain of SA empty
 - \Rightarrow Backtrack



	WA	NT	Q	NSW	V	SA	T
Initial domains	[red, green, blue]						
After WA=red	[red]	[green, blue]	[red, green, blue]	[red, green, blue]	[red, green, blue]	[green, blue]	[red, green, blue]
After Q=green	[red]	[blue]	[green]	[red, blue]	[red, green, blue]	[blue]	[red, green, blue]
After V=blue	[red]	[blue]	[green]	[red]	[blue]		[red, green, blue]

Interleaved Search and Inference

- Combining MRV with forward checking is more effective
 - After assigning {WA=red} NT and SA have two values.
 - MRV would have chosen one of them first leading to a solution
- Forward checking incrementally computes the information that the MRV heuristic needs...



	WA	NT	Q	NSW	V	SA	T
Initial domains	[Red, Green, Blue]						
After WA=red	[Red]	[Green, Blue]	[Red, Green, Blue]	[Red, Green, Blue]	[Red, Green, Blue]	[Green, Blue]	[Red, Green, Blue]
After Q=green	[Red]	[Blue]	[Green]	[Red, Blue]	[Red, Green, Blue]	[Blue]	[Red, Green, Blue]
After V=blue	[Red]	[Blue]	[Green]	[Red]	[Blue]		[Red, Green, Blue]

Interleaved Search and Inference

- Forward checking doesn't detect all inconsistencies since it does not look ahead far enough
 - In the $Q=\text{green}$ row, both NT and SA are left with blue as their only possible value, which is an inconsistency, since they are neighbors.



	WA	NT	Q	NSW	V	SA	T
Initial domains	[Red, Green, Blue]						
After $WA=\text{red}$	[Red]	[Green, Blue]	[Red, Green, Blue]	[Red, Green, Blue]	[Red, Green, Blue]	[Green, Blue]	[Red, Green, Blue]
After $Q=\text{green}$	[Red]	[Blue]	[Green]	[Red]	[Red, Green, Blue]	[Blue]	[Red, Green, Blue]
After $V=\text{blue}$	[Red]	[Blue]	[Green]	[Red]	[Blue]		[Red, Green, Blue]

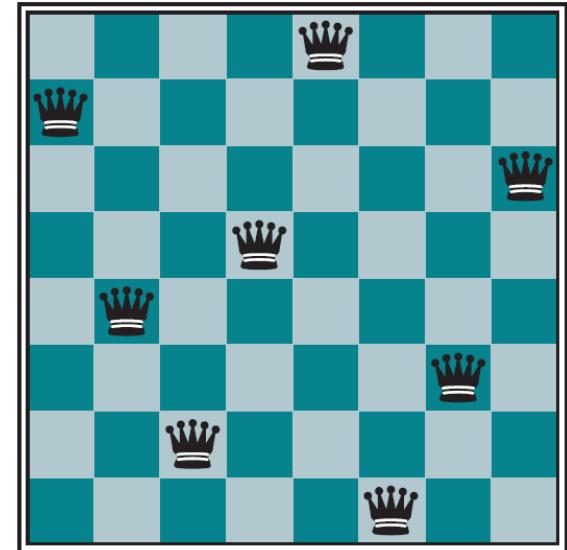
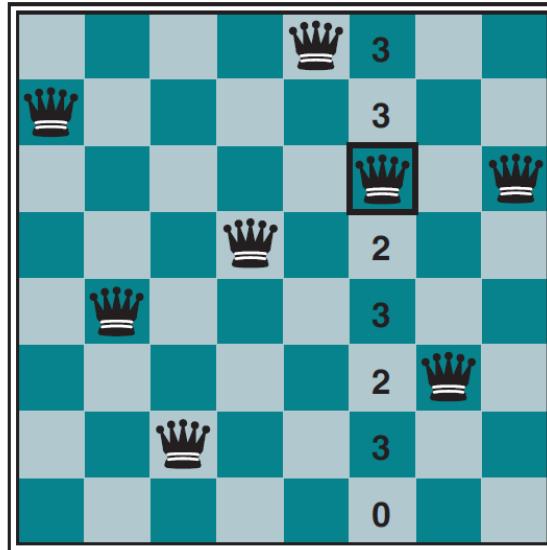
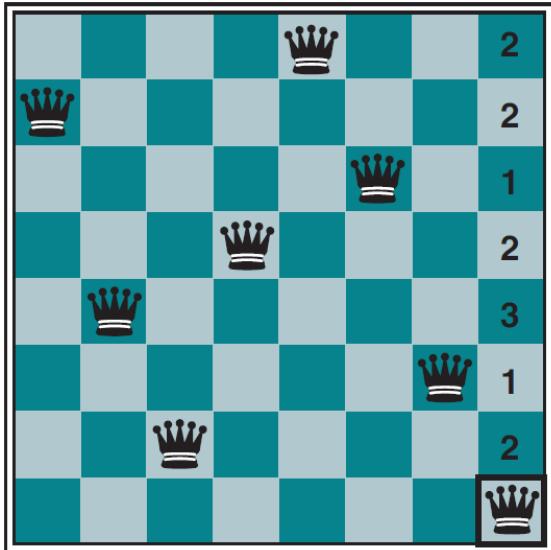
Maintaining Arc Consistency (MAC)

- After a variable X_i is assigned a value, inference calls AC-3
 - Instead of a queue of all the arcs, it starts with only the arcs (X_j, X_i) for all X_j that are unassigned and are neighbors of X_i .
 - If any variable's domain is reduced to the empty set, the call to AC-3 fails which triggers backtracking immediately.
- We can see that MAC is strictly more powerful than forward checking.
 - Unlike MAC, forward checking does not recursively propagate constraints.

Local Search

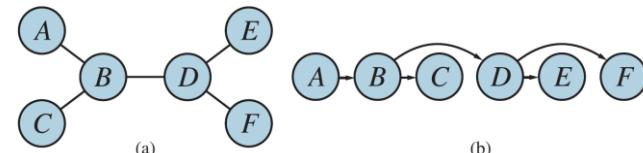
- Start with a possibly conflicting but complete assignment
- Pick a conflicted variable randomly and change its value
- Use Min-Conflicts heuristic to find a new value
 - Select the value that results in the smallest number of conflicts
 - Quite effective. Can solve a million variable N-Queens in 50 or so steps!
- Local search is great in online settings since repairing is usually much faster than solving from scratch

Local Search



Tree Structured CSP

- CSP is NP-complete in general BUT any tree structured CSP can be solved in linear time in the number of CSP variables
- Directional Arc Consistency (DAC)
 - Given an ordering of variables X_1, X_2, \dots, X_n , a CSP is DAC iff every X_i is arc-consistent with X_j where $j > i$
 - Linear time algorithm: topological sort the variables, make the tree DAC and traverse from root down the leaves picking any remaining values.



Quiz

1. What is the primary purpose of arc consistency (AC-3) in a CSP solver?
 - A. To identify a single correct value for each variable.
 - B. To remove values from a variable's domain that cannot possibly be part of a consistent solution.
 - C. To add new constraints to the problem to make it easier to solve.
2. When a search algorithm for a CSP finds a partial assignment that violates a constraint, what does it do?
 - A. It changes the domain of the conflicted variable to resolve the issue.
 - B. It adds a new variable to the problem.
 - C. It backtracks to the most recent variable with remaining values in its domain.

Quiz

1. What is the primary purpose of arc consistency (AC-3) in a CSP solver?
 - A. To identify a single correct value for each variable.
 - B. To remove values from a variable's domain that cannot possibly be part of a consistent solution.
 - C. To add new constraints to the problem to make it easier to solve.
2. When a search algorithm for a CSP finds a partial assignment that violates a constraint, what does it do?
 - A. It changes the domain of the conflicted variable to resolve the issue.
 - B. It adds a new variable to the problem.
 - C. It backtracks to the most recent variable with remaining values in its domain.

Quiz

3. What is the purpose of the **least constraining value** heuristic when solving a CSP?
 - A. To choose the value that has the fewest constraints associated with it.
 - B. To choose the value that is most likely to lead to a solution.
 - C. To choose the value that prunes the smallest number of values from the domains of neighboring variables.
4. What is the key difference between backtracking search and local search for CSPs?
 - A. Backtracking uses a single variable, while local search considers multiple variables at once.
 - B. Backtracking is a DFS, while local search starts with a full assignment and improves it.
 - C. Backtracking is incomplete, while local search is complete.
5. What is the min-conflicts heuristic used for in local search for CSPs?
 - A. To select the variable that is in a conflict.
 - B. To choose which variable to assign a value to next.
 - C. To choose a value for a conflicted variable that results in the minimum number of conflicts with other variables.

Quiz

3. What is the purpose of the **least constraining value** heuristic when solving a CSP?
 - A. To choose the value that has the fewest constraints associated with it.
 - B. To choose the value that is most likely to lead to a solution.
 - C. **To choose the value that prunes the smallest number of values from the domains of neighboring variables.**

3. What is the key difference between backtracking search and local search for CSPs?
 - A. Backtracking uses a single variable, while local search considers multiple variables at once.
 - B. **Backtracking is a DFS, while local search starts with a full assignment and improves it.**
 - C. Backtracking is incomplete, while local search is complete.

4. What is the min-conflicts heuristic used for in local search for CSPs?
 - A. To select the variable that is in a conflict.
 - B. To choose which variable to assign a value to next.
 - C. **To choose a value for a conflicted variable that results in the minimum number of conflicts with other variables.**

Summary

- A CSP is defined by three components:
 - A set of variables
 - A domain of possible values for each variable
 - A set of constraints that specify which combinations of values are allowed.
- Constraints can be classified by the number of variables they involve.
 - A unary constraint affects a single variable, a binary constraint involves two variables, and a global constraint affects more than two.
- Many real-world problems can be effectively formulated as CSPs

Summary

- **Backtracking Search** is a core algorithm for solving CSPs.
 - It works by performing a depth-first search on the variables. The algorithm incrementally assigns a value to one variable at a time, and if a variable assignment leads to a constraint violation, it backtracks to the previous variable and tries a different value.
- **Heuristics** are used to improve backtracking search
 - **Variable Ordering:** The **Minimum Remaining Values** heuristic selects the variable with the fewest legal values remaining.
 - **Value Ordering:** The **least constraining value heuristic** selects the value that rules out the fewest choices for neighboring variables.

Summary

- **Constraint Propagation** is a technique used to prune the search space. It works by using constraints to reduce the domain of variables.
 - **Arc consistency (AC-3)** is a popular algorithm that removes values from a variable's domain that have no consistent value in a neighboring variable's domain.
- **Forward Checking** checks for conflicts with unassigned variables immediately after a variable is assigned.
 - Prunes the domains of neighboring unassigned variables, preventing future dead ends.

Summary

- Local Search algorithms start with a complete but possibly invalid assignment and iteratively improve it by making small changes.
 - Min-Conflicts Heuristic guides the iterative repair process in local search.
 - For a variable is in conflict, it chooses a new value for that variable that minimizes the number of remaining conflicts with other variables.
- Completeness vs. Efficiency: Backtracking search is a complete algorithm, and local search algorithms, while often faster, are typically incomplete.
 - Local search may get stuck in a local optimum and fail to find a solution.

- **Reading:** Chapter 6
- **Assignments:** PS 4, csp.ipynb
- **Next:** Logical Agents, Chapter 7