

# Artificial Intelligence

## 6. CSP

Shashi Prabh

School of Engineering and Applied Science  
Ahmedabad University

# Contents

Goal: use factored representation of agents to solve problems.

## Topics

- Constraint Satisfaction Problem
- Constraint Propagation
- Backtracking Search
- Local Search

# Constraint Satisfaction Problems (CSP)

- We consider factored representation of states
  - A state is a set of variables
  - A problem solution is an assignment of values to the state variables where all the constraints on the variables are satisfied

# Constraint Satisfaction Problems (CSP)

- Why CSP?
  - CSP is a natural formulation in many problems
    - Scheduling, planning, resource allocation, temporal models, control etc.
  - Significant reduction of search space, availability of fast solvers
  - Insight into the problem structure can be used for search speed-up
    - Some intractable atomic search-space problems can be quickly solved as CSP formulation
  - Actions and transition model can be deduced from the formulation

# Constraint Satisfaction Problems (CSP)

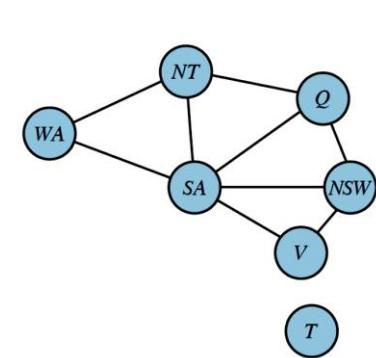
- A CSP consists of three components ( $X$ ,  $D$ ,  $C$ ):
- **Variables**  $X = \{x_1, x_2, \dots, x_n\}$
- **Domains**  $D = \{D_1, D_2, \dots, D_n\}$
- **Constraints**  $C = \{c_1, c_2, \dots, c_m\}$ 
  - Domain  $D_i$  consists of the set of **allowable values**  $\{v_1, \dots, v_k\}$  for each  $x_i$ 
    - $\{T, F\}$  for a Boolean variable
  - Constraint  $c_j$  consists of a pair  **$\langle$ scope, relation $\rangle$** 
    - $\langle(x_1, x_2), x_1 \neq x_2 \rangle$  or just  $x_1 \neq x_2$

# Constraint Satisfaction Problems (CSP)

- Goal: Assign values to the variables from their respective domains such that all the constraints are satisfied
  - An assignment that does not violate any constraint is called **consistent** or **legal assignment**
  - A solution to a CSP is a complete and consistent assignment
  - Solving a CSP is NP-complete in general
- Types of constraints:
  - Unary, binary, global
  - $\langle (x_1), x_1 \neq 4 \rangle$ ,  $\langle (x_1, x_2), x_1 \neq x_2 \rangle$ , AllDiff, AtMost, AtLeast

# Map coloring

- $X = \{W, N, S, Q, NSW, V, T\}$
- $D = \{r, g, b\}$
- $C = \{ W \neq N, S \neq N, Q \neq N, W \neq S, S \neq Q, \text{etc} \}$ 
  - $W \neq N$  means  $\{(r, g), (r, g), (g, r), (g, b), (b, r), (b, g)\}$
  - Note the reduced search space due to the constraints:  $2^5$  vs  $3^5$
  - **Can you find one solution?**
- In a CSP **constraint graph**, two variables are connected by an edge if there is a constraint that involves both



# Job-Shop Scheduling – Car Assembly

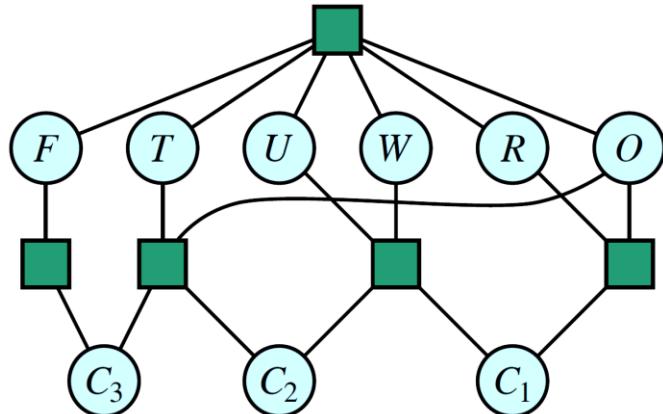
- $X$  is the set of tasks

$\{Axe_F, Axe_B, Wheel_{RF}, Wheel_{LF}, Wheel_{RB}, Wheel_{LB}, Nuts_{RF}, Nuts_{LF}, Nuts_{RB}, Nuts_{LB}, Cap_{RF}, Cap_{LF}, Cap_{RB}, Cap_{LB}, Inspect\}$

- Values are the start times of tasks:  $D_i = \{0, 1, \dots, 30\}$
- Constraints: precedence constraints, completion times
  - If  $T_i$  precedes  $T_j$ ,  $T_i + D_i \leq T_j$ 
$$Axe_F + 10 \leq Wheel_{RF}; \quad Axe_F + 10 \leq Wheel_{LF}; \\ Axe_B + 10 \leq Wheel_{RB}; \quad Axe_B + 10 \leq Wheel_{LB}.$$
  - Disjunctions: Axle installations must not overlap in time
$$(Axe_F + 10 \leq Axe_B) \quad \text{or} \quad (Axe_B + 10 \leq Axe_F)$$
  - Exercise: CSP formulation of 8-Queens problem

# Cryptarithmetic Puzzles

$$\begin{array}{r} T \quad W \quad O \\ + \quad T \quad W \quad O \\ \hline F \quad O \quad U \quad R \end{array}$$



Constraint hypergraph

- Constraints: AllDiff (F, T, U, W, R, O), F ≠ 0 and

$$O + O = R + 10 \cdot C_1$$

$$C_1 + W + W = U + 10 \cdot C_2$$

$$C_2 + T + T = O + 10 \cdot C_3$$

$$C_3 = F,$$

1. In the context of a Constraint Satisfaction Problem (CSP), what are the three main components?
  - A. Variables, Functions, and Relations
  - B. Nodes, Arcs, and Paths
  - C. States, Transitions, and Goals
  - D. Variables, Domains, and Constraints
2. Which of the following is an example of a binary constraint?
  - A. The sum of variables A, B, and C must be less than 10.
  - B. The value of variable A must be different from the value of variable B.
  - C. The colors of all bordering regions must be different.
  - D. The value of variable A must be even.

1. In the context of a Constraint Satisfaction Problem (CSP), what are the three main components?
  - A. Variables, Functions, and Relations
  - B. Nodes, Arcs, and Paths
  - C. States, Transitions, and Goals
  - D. Variables, Domains, and Constraints
2. Which of the following is an example of a binary constraint?
  - A. The sum of variables A, B, and C must be less than 10.
  - B. The value of variable A must be different from the value of variable B.
  - C. The colors of all bordering regions must be different.
  - D. The value of variable A must be even.

3. In the N-queens problem, what do the variables, domains, and constraints represent?
- A. Variables are the queens, domains are the squares on the board, and constraints are that no two queens can attack each other.
  - B. Variables are the rows and columns, domains are whether a queen is present, and constraints are that a queen must be in every row.
  - C. Variables are the rows, domains are the columns, and constraints are that queens cannot be in the same row.
  - D. Variables are the columns, domains are the rows, and constraints are that no two queens can share a row, column, or diagonal.
4. Which of the following problems can be formulated as a CSP?
- A. Generating a creative story given a set of keywords.
  - B. Calculating the final grade for a student.
  - C. Determining a weekly schedule for classes without any conflicts.
  - D. Finding the shortest path between two cities on a map.

1. In the N-queens problem, what do the variables, domains, and constraints represent?
  - A. Variables are the queens, domains are the squares on the board, and constraints are that no two queens can attack each other.
  - B. Variables are the rows and columns, domains are whether a queen is present, and constraints are that a queen must be in every row.
  - C. Variables are the rows, domains are the columns, and constraints are that queens cannot be in the same row.
  - D. Variables are the columns, domains are the rows, and constraints are that no two queens can share a row, column, or diagonal.
2. Which of the following problems can be formulated as a CSP?
  - A. Generating a creative story given a set of keywords.
  - B. Calculating the final grade for a student.
  - C. Determining a weekly schedule for classes without any conflicts.
  - D. Finding the shortest path between two cities on a map.

# Constraint Propagation

- A CSP algorithm can generate successors as new assignments
- Constraint propagation is an alternative where using constraints the number of legal values is reduced
- Used along-with search and/or as a preprocessing step

# Constraint Propagation

- Local consistency shrinks the search space by eliminating the inconsistent assignments
- Types of local consistency
  - Node consistency
  - Arc consistency
    - Path and K-Consistency
- Global constraints, bounds propagation

# Node Consistency

- A node in the constraint graph is node-consistent if all the values in the variable's domain satisfy the variable's unary constraints.
- Example: consider a unary constraint  $SA \neq \{\text{green}\}$ 
  - The variable SA with initial domain {red, green, blue} can be made node consistent by eliminating green from its domain, leaving SA with the reduced domain {red, blue}.
- A graph is node-consistent if every variable in the graph is node-consistent.
  - One can just eliminate domain values inconsistent with unary constraints.

# Arc Consistency

- A variable is arc-consistent if for every value in its domain, there is some value in the domains of all the variables connected by a binary constraint.
  - Example: consider the constraint  $Y = X^2$ ,  $D_X = \mathbb{N}$ ,  $D_Y = \{ 1, 4, 9\}$
  - $X$  is made arc-consistent with  $Y$  by restricting  $D_X = \{ 1, 2, 3\}$
  - However, arc-consistency is ineffective in the map coloring example
- **AC-3** is a widely used arc-consistency algorithm

# AC-3 (Mackworth, 1977)

**function** AC-3(*csp*) **returns** false if an inconsistency is found and true otherwise  
*queue*  $\leftarrow$  a queue of arcs, initially all the arcs in *csp*

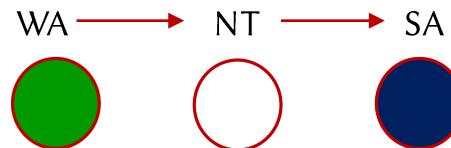
**while** *queue* is not empty **do**  
     $(X_i, X_j) \leftarrow \text{POP}(\text{queue})$   
    **if** REVISE(*csp*, *X<sub>i</sub>*, *X<sub>j</sub>*) **then**  
        **if** size of *D<sub>i</sub>* = 0 **then return** false  
        **for each** *X<sub>k</sub>* **in** *X<sub>i</sub>.NEIGHBORS - {X<sub>j</sub>}* **do**  
            add (*X<sub>k</sub>*, *X<sub>i</sub>*) to *queue*  
**return** true

**function** REVISE(*csp*, *X<sub>i</sub>*, *X<sub>j</sub>*) **returns** true iff we revise the domain of *X<sub>i</sub>*  
*revised*  $\leftarrow$  false  
**for each** *x* **in** *D<sub>i</sub>* **do**  
    **if** no value *y* in *D<sub>j</sub>* allows (*x,y*) to satisfy the constraint between *X<sub>i</sub>* and *X<sub>j</sub>* **then**  
        delete *x* from *D<sub>i</sub>*  
        *revised*  $\leftarrow$  true  
**return** *revised*

- Initially, each binary constraint inserts two arcs
- $X_i$  is being made consistent with  $X_j$
- $O(c d^3)$  worst-case complexity

# Path Consistency

- AC does not help with map coloring
  - Does not object to 2-coloring the map
- $\{X_i, X_j\}$  is path-consistent with respect to a third variable  $X_m$  if, for every assignment  $\{X_i = a, X_j = b\}$  consistent with the constraints on  $\{X_i, X_j\}$ , there is an assignment to  $X_m$  that satisfies the constraints on  $\{X_i, X_m\}$  and  $\{X_m, X_j\}$ .
  - Refers to the overall consistency of the path  $X_i \rightarrow X_m \rightarrow X_j$
  - Can infer that no valid 2-coloring of the Australia map exists



# K-Consistency

- A CSP is **k-consistent** if, for any set of **k-1** variables and for any consistent assignment to those variables, a consistent value can always be assigned to any  $k^{\text{th}}$  variable
  - 1-consistency says that, given the empty set, we can make any set of one variable consistent: this is what we called node consistency
  - 2-consistency is the same as arc consistency
  - 3-consistency (binary constraints) is the same as path consistency

# K-Consistency

- A CSP is **strongly k-consistent** if it is k-consistent and is also  $(k-1)$ -consistent,  $(k-2)$ -consistent, . . . , 1-consistent
- Why?
  - Can design a greedy algorithm
  - **CSP is NP-complete**
    - K-consistency requires exponential time and space

# Global constraints

- A global constraint involves an arbitrary number of variables. It is more efficient to handle these by special-purpose algorithms
  - **AllDiff**: if  $m$  variables are involved in an AllDiff constraint, and if  $n$  possible distinct values altogether are available, then the constraint cannot be satisfied if  $m > n$
  - **Atmost**: resource constraint
    - Example: no more than 10 personnel are scheduled in total
    - We can detect an inconsistency simply by checking the sum of the **minimum** values of the current domains

# Global constraints

- **Bounds propagation:** For problems with large integer domains it is usually not efficient to represent the domain of each variable as a large set of integers.
  - Domains can be represented by upper and lower bounds and managed by bounds propagation

# Global constraints

- Example:
  - Consider two flights, F1 and F2, for which the planes have capacities 165 and 385, respectively
  - The initial domains for the numbers of passengers are then  $D_1 = [0, 165]$  and  $D_2 = [0, 385]$
  - The additional constraint that **the two flights together must carry 450 people** can be handled by propagating bounds constraints as  $D_1 = [65, 165]$  and  $D_2 = [285, 385]$

# Sudoku

	1	2	3	4	5	6	7	8	9
A			3	2		6			
B	9		3		5				1
C			1	8		6	4		
D			8	1		2	9		
E	7								8
F			6	7		8	2		
G			2	6		9	5		
H	8		2		3				9
I			5		1	3			

	1	2	3	4	5	6	7	8	9
A	4	8	3	9	2	1	6	5	7
B	9	6	7	3	4	5	8	2	1
C	2	5	1	8	7	6	4	9	3
D	5	4	8	1	3	2	9	7	6
E	7	2	9	5	6	4	1	3	8
F	1	3	6	7	9	8	2	4	5
G	3	7	2	6	8	9	5	1	4
H	8	1	4	2	5	3	7	6	9
I	6	9	5	4	1	7	3	8	2

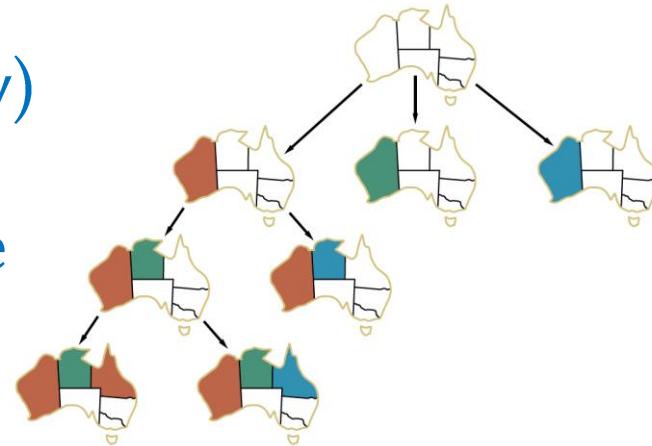
Exercise: Write CSP formulation!

# Backtracking Search

- DFS at heart but with significant differences
- Search for solution is needed when after constraint propagation there exist variables with  $\geq 1$  values
- For a CSP with  $n$  variables of domain size  $d$ , a pure DFS results in a search tree with  $n! d^n$  leaf nodes at depth  $n$ 
  - The branching factor at the top would be  $nd$ , at the next level  $(n-1)d$  and so on
- The order of assignments does not matter
  - There are only  $d^n$  possible assignments!

# Backtracking Search

- Backtracking search progresses via a recursive call
- An unassigned variable is (repeatedly) chosen, a value is assigned and the search progresses to another variable
  - If the search succeeds, the solution is returned
  - If the search fails, the assignment is restored to the previous state, and the next value is tried



# Backtracking Search

```
function BACKTRACKING-SEARCH(csp) returns a solution or failure
    return BACKTRACK(csp, {})

function BACKTRACK(csp, assignment) returns a solution or failure
    if assignment is complete then return assignment
    var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(csp, assignment)
    for each value in ORDER-DOMAIN-VALUES(csp, var, assignment) do
        if value is consistent with assignment then
            add  $\{var = value\}$  to assignment
            inferences  $\leftarrow$  INFERENCE(csp, var, assignment)
            if inferences  $\neq$  failure then
                add inferences to csp
                result  $\leftarrow$  BACKTRACK(csp, assignment)
                if result  $\neq$  failure then return result
                remove inferences from csp
                remove  $\{var = value\}$  from assignment
    return failure
```

# Improving Backtracking Search

- Backtracking search can be improved using **domain-independent heuristics** that take advantage of the factored representation of states
- Variable and value ordering heuristics
  - **Minimum-remaining-values heuristic (MRV)**
    - Start with F in cryptarithmetric puzzle
  - **Degree heuristic – largest first**
    - Start with SA in Australia map
  - **Least constraining value first heuristic (LCV)**
    - Values that rule out the fewest choices first

# Forward Checking

- **Forward Checking:** Check for arc consistency upon a variable assignment
  - Upon assignment to X, make each unassigned variable Y that is connected to X by a constraint, arc-consistent with X
    - After assigning V =blue, forward checking finds domain of SA empty
      - $\Rightarrow$  Backtrack



	WA	NT	Q	NSW	V	SA	T
Initial domains	[Red, Green, Blue]						
After WA=red	[Red]	[Green, Blue]	[Red, Green, Blue]	[Red, Green, Blue]	[Red, Green, Blue]	[Green, Blue]	[Red, Green, Blue]
After Q=green	[Red]	[Blue]	[Green]	[Red, Blue]	[Red, Green, Blue]	[Blue]	[Red, Green, Blue]
After V=blue	[Red]	[Blue]	[Green]	[Red]	[Blue]		[Red, Green, Blue]

# Interleaved Search and Inference

- Combining MRV with forward checking is more effective
  - After assigning {WA=red} NT and SA have two values.
  - MRV would have chosen one of them first leading to a solution
- Forward checking incrementally computes the information that the MRV heuristic needs...



	WA	NT	Q	NSW	V	SA	T
Initial domains	[Red, Green, Blue]						
After WA=red	[Red]	[Green, Blue]	[Red, Green, Blue]	[Red, Green, Blue]	[Red, Green, Blue]	[Green, Blue]	[Red, Green, Blue]
After Q=green	[Red]	[Blue]	[Green]	[Red, Blue]	[Red, Green, Blue]	[Blue]	[Red, Green, Blue]
After V=blue	[Red]	[Blue]	[Green]	[Red]	[Blue]		[Red, Green, Blue]

# Interleaved Search and Inference

- Forward checking doesn't detect all inconsistencies since it does not look ahead far enough
  - In the  $Q=\text{green}$  row, both NT and SA are left with blue as their only possible value, which is an inconsistency, since they are neighbors.



	WA	NT	Q	NSW	V	SA	T
Initial domains	[Red, Green, Blue]						
After $WA=\text{red}$	[Red]	[Green, Blue]	[Red, Green, Blue]	[Red, Green, Blue]	[Red, Green, Blue]	[Green, Blue]	[Red, Green, Blue]
After $Q=\text{green}$	[Red]	[Blue]	[Green]	[Red]	[Red, Green, Blue]	[Blue]	[Red, Green, Blue]
After $V=\text{blue}$	[Red]	[Blue]	[Green]	[Red]	[Blue]		[Red, Green, Blue]

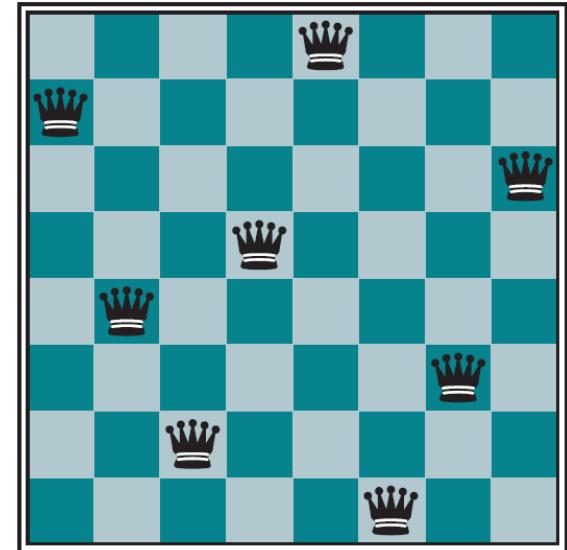
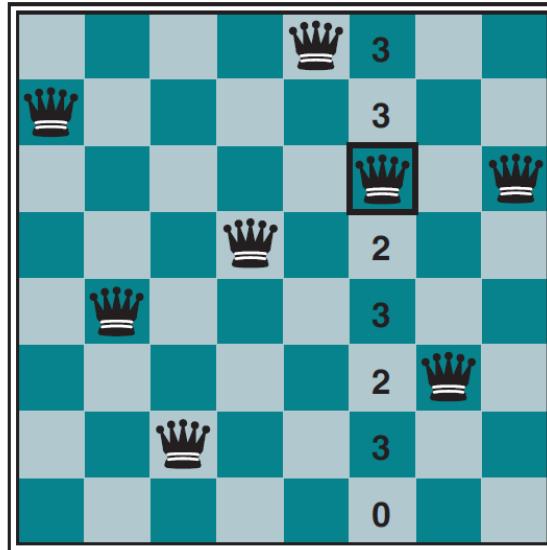
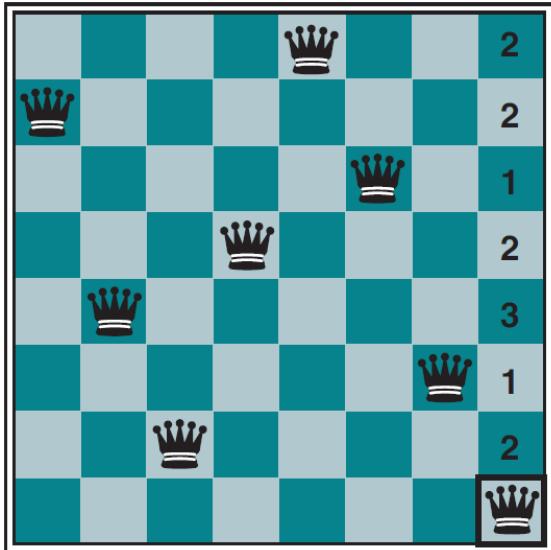
# Maintaining Arc Consistency (MAC)

- After a variable  $X_i$  is assigned a value, inference calls AC-3
  - Instead of a queue of all the arcs, it starts with only the arcs  $(X_j, X_i)$  for all  $X_j$  that are unassigned and are neighbors of  $X_i$ .
  - If any variable's domain is reduced to the empty set, the call to AC-3 fails which triggers backtracking immediately.
- We can see that MAC is strictly more powerful than forward checking.
  - Unlike MAC, forward checking does not recursively propagate constraints.

# Local Search

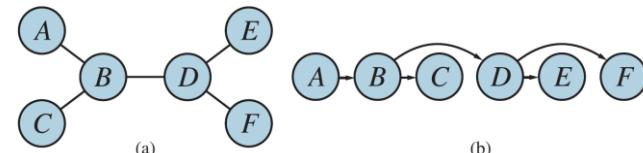
- Start with a possibly conflicting but complete assignment
- Pick a conflicted variable randomly and change its value
- Use Min-Conflicts heuristic to find a new value
  - Select the value that results in the smallest number of conflicts
  - Quite effective. Can solve a million variable N-Queens in 50 or so steps!
- Local search is great in online settings since repairing is usually much faster than solving from scratch

# Local Search



# Tree Structured CSP

- CSP is NP-complete in general BUT any tree structured CSP can be solved in linear time in the number of CSP variables
- Directional Arc Consistency (DAC)
  - Given an ordering of variables  $X_1, X_2, \dots, X_n$ , a CSP is DAC iff every  $X_i$  is arc-consistent with  $X_j$  where  $j > i$
  - Linear time algorithm: topological sort the variables, make the tree DAC and traverse from root down the leaves picking any remaining values.



# Quiz

1. What is the primary purpose of arc consistency (AC-3) in a CSP solver?
  - A. To identify a single correct value for each variable.
  - B. To remove values from a variable's domain that cannot possibly be part of a consistent solution.
  - C. To add new constraints to the problem to make it easier to solve.
2. When a search algorithm for a CSP finds a partial assignment that violates a constraint, what does it do?
  - A. It changes the domain of the conflicted variable to resolve the issue.
  - B. It adds a new variable to the problem.
  - C. It backtracks to the most recent variable with remaining values in its domain.

# Quiz

1. What is the primary purpose of arc consistency (AC-3) in a CSP solver?
  - A. To identify a single correct value for each variable.
  - B. To remove values from a variable's domain that cannot possibly be part of a consistent solution.
  - C. To add new constraints to the problem to make it easier to solve.
2. When a search algorithm for a CSP finds a partial assignment that violates a constraint, what does it do?
  - A. It changes the domain of the conflicted variable to resolve the issue.
  - B. It adds a new variable to the problem.
  - C. It backtracks to the most recent variable with remaining values in its domain.

# Quiz

3. What is the purpose of the **least constraining value** heuristic when solving a CSP?
  - A. To choose the value that has the fewest constraints associated with it.
  - B. To choose the value that is most likely to lead to a solution.
  - C. To choose the value that prunes the smallest number of values from the domains of neighboring variables.
4. What is the key difference between backtracking search and local search for CSPs?
  - A. Backtracking uses a single variable, while local search considers multiple variables at once.
  - B. Backtracking is a DFS, while local search starts with a full assignment and improves it.
  - C. Backtracking is incomplete, while local search is complete.
5. What is the min-conflicts heuristic used for in local search for CSPs?
  - A. To select the variable that is in a conflict.
  - B. To choose which variable to assign a value to next.
  - C. To choose a value for a conflicted variable that results in the minimum number of conflicts with other variables.

# Quiz

3. What is the purpose of the **least constraining value** heuristic when solving a CSP?
  - A. To choose the value that has the fewest constraints associated with it.
  - B. To choose the value that is most likely to lead to a solution.
  - C. **To choose the value that prunes the smallest number of values from the domains of neighboring variables.**
  
3. What is the key difference between backtracking search and local search for CSPs?
  - A. Backtracking uses a single variable, while local search considers multiple variables at once.
  - B. **Backtracking is a DFS, while local search starts with a full assignment and improves it.**
  - C. Backtracking is incomplete, while local search is complete.
  
4. What is the min-conflicts heuristic used for in local search for CSPs?
  - A. To select the variable that is in a conflict.
  - B. To choose which variable to assign a value to next.
  - C. **To choose a value for a conflicted variable that results in the minimum number of conflicts with other variables.**

# Summary

- A CSP is defined by three components:
  - A set of variables
  - A domain of possible values for each variable
  - A set of constraints that specify which combinations of values are allowed.
- Constraints can be classified by the number of variables they involve.
  - A unary constraint affects a single variable, a binary constraint involves two variables, and a global constraint affects more than two.
- Many real-world problems can be effectively formulated as CSPs

# Summary

- **Backtracking Search** is a core algorithm for solving CSPs.
  - It works by performing a depth-first search on the variables. The algorithm incrementally assigns a value to one variable at a time, and if a variable assignment leads to a constraint violation, it backtracks to the previous variable and tries a different value.
- **Heuristics** are used to improve backtracking search
  - **Variable Ordering:** The **Minimum Remaining Values** heuristic selects the variable with the fewest legal values remaining.
  - **Value Ordering:** The **least constraining value heuristic** selects the value that rules out the fewest choices for neighboring variables.

# Summary

- **Constraint Propagation** is a technique used to prune the search space. It works by using constraints to reduce the domain of variables.
  - **Arc consistency (AC-3)** is a popular algorithm that removes values from a variable's domain that have no consistent value in a neighboring variable's domain.
- **Forward Checking** checks for conflicts with unassigned variables immediately after a variable is assigned.
  - Prunes the domains of neighboring unassigned variables, preventing future dead ends.

# Summary

- Local Search algorithms start with a complete but possibly invalid assignment and iteratively improve it by making small changes.
  - Min-Conflicts Heuristic guides the iterative repair process in local search.
    - For a variable is in conflict, it chooses a new value for that variable that minimizes the number of remaining conflicts with other variables.
- Completeness vs. Efficiency: Backtracking search is a complete algorithm, and local search algorithms, while often faster, are typically incomplete.
  - Local search may get stuck in a local optimum and fail to find a solution.

- **Reading:** Chapter 6
- **Assignments:** PS 4, csp.ipynb
- **Next:** Logical Agents, Chapter 7